

Name 1:

Name 2:

COMPUTER NETWORKING

LAB EXERCISES (TP) 2

L1 v.s. L2 v.s. L3, NAT, PHYSICAL CONNECTION, AND TROUBLESHOOTING

October 14, 2016

Abstract

In this Lab you will work with the virtual environment introduced in Lab 1 as well as a physical machine in the INF019. First you will see the different behavior of networking devices that work on Layer 1, layer 2 and layer 3; then you will configure your virtual network to be able to access the Internet; later you will connect your virtual network to the Internet using a machine in the INF 019; and finally, you will help Jon and Arya to fix their networking problems when a common enemy, Joffrey, changes the configuration in their network.

1 PREPARING THE LAB

1.1 LAB REPORT

Type your answers in this document. We recommend you use Adobe Reader XI to open this PDF. When you finish, save the report and upload it on moodle. Don't forget to write your names on the first page of the report. **The deadline is Wednesday, October 25, 23:59:59**

1.2 SETTING UP VIRTUAL MACHINE

In this Lab, you will work with the same virtual machine that you created in Lab 1. Copy the **lab2 resources** folder from Moodle into the shared folder of your VM before starting the lab.

1.3 USING SCRIPTS

As a general advice, use scripts to save your work for each section, especially Section 3. This is useful for 1) saving time and not repeating the same commands each time you restart Mininet, and 2) reviewing and debugging your work in case you run into issues.

2 LAYER 1 VS. LAYER 2 VS. LAYER 3 NETWORKING

The aim of this section is to illustrate the difference between networking devices that work at layer 1, layer 2 and layer 3. For this exercise we only consider IPv4 addressing.

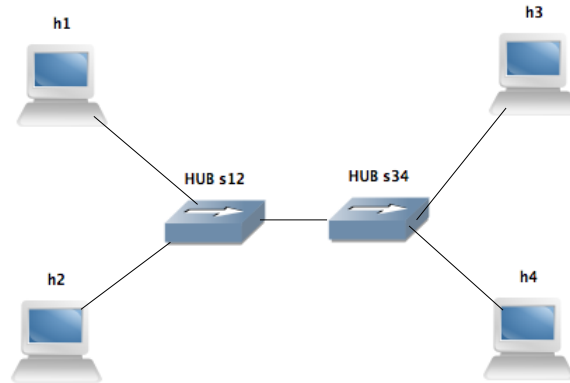


Figure 1: Network configuration with two hubs in the center

2.1 USING HUB AS A NETWORKING DEVICE

A hub is layer-one intermediate system that repeats bits (symbols) received from one of its ports to all other ports. In this section we analyze how it works.

Open a terminal in your VM and run the script `topo1.py`, which should be located in the shared folder on the Desktop. If not, refer to Section 1.2.

```
# python topo1.py
```

This will create the network described in Figure 1, and redirect you to the mininet CLI. Additionally, one terminal will appear for each of the four hosts. The four new terminals will be labeled (h1, h2, h3, h4) for convenience.

Run the following two commands to configure s12 and s34 as hubs.

```
mininet> sh ovs-ofctl add-flow s12 action=flood
mininet> sh ovs-ofctl add-flow s34 action=flood
```

h1, h2, and h3 should be located on the 10.0.0.0/24 subnet with the fourth byte of their IP address being 1, 2, and 3, respectively. h4 should have the IP address 10.0.1.4/24. Check the configuration of the IP addresses for each of the hosts in their respective terminals and correct any wrong configuration.



Q1/ Was there a wrongly configured host? How did you find out?

[A1]

Q2/ Which line of code in the mininet configuration script does this error correspond to?

[A2]

Q3/ What command did you run to fix the configuration issue?

[A3]

Start Wireshark on all four hosts. It will be hard to keep track of which Wireshark window corresponds to which host. One way to do so would be to start Wireshark on the hosts in order, i.e. h1, then h2, then h3, and finally on h4. This way the Wireshark windows will be in this same order in the taskbar. Start capturing on all the eth0 interfaces.

```
# wireshark &
```

From h1, ping h2.

```
mininet> h1 ping h2
```



Q4/ What would be another way of doing the same thing?

[A4]

Q5/ Is there any difference between the traffic captured by the four hosts in Wireshark?

[A5]

Q6/ Explain why you see these results.

[A6]

2.2 USING A BRIDGE AS A NETWORKING DEVICE

A bridge is a link-layer intermediate system which expands a LAN by making forwarding decision based on destination MAC-address. In this section you will learn how they work.

We will first change s34 to act as a bridge by running the following command.

```
mininet> sh ovs-ofctl add-flow s34 action=normal
```

Note that `s12` and `s34` were originally created as bridges, but we modified them in the previous section to act as hubs by adding flows. You will learn more about flows in Lab 4.



Q7/ If you were to exit Mininet and run `topo1.py` again, what command(s) would you then execute to reach the topology we have now?

[A7]

Now, let's test our bridge configuration. Start a Wireshark capture on all four hosts again, and again ping `h2` from `h1`.



Q8/ Describe the different types of packets observed on `h1`, `h2`, `h3` and `h4`.

[A8]

Q9/ Explain the results. What is the difference compared to having two hubs?

[A9]

Q10/ Ping from `h1` to `h4`. Observe the traffic captured and explain your findings.

[A10]

Now, focus on the interfaces of `h1`, `h2` and `h3`. Ping from `h1` to `h3`.



Q11/ Observe the packets captured on `h2` and explain the results.

[A11]

Q12/ Compare the packets sent by `h1` to the ones received by `h3`, specifically at source/destination MAC-addresses. Explain the similarities and differences, if any.

[A12]

2.3 USING A ROUTER AS A NETWORKING DEVICE

We have already configured a router in Lab 1, but we did not address how it worked. In this section we learn about the process of routing a packet.

We'll begin by exiting Mininet, cleaning up the previous topology, and running `topo2.py`. The new topology consists of three hosts, a router, and a bridge, as shown in Figure 2.

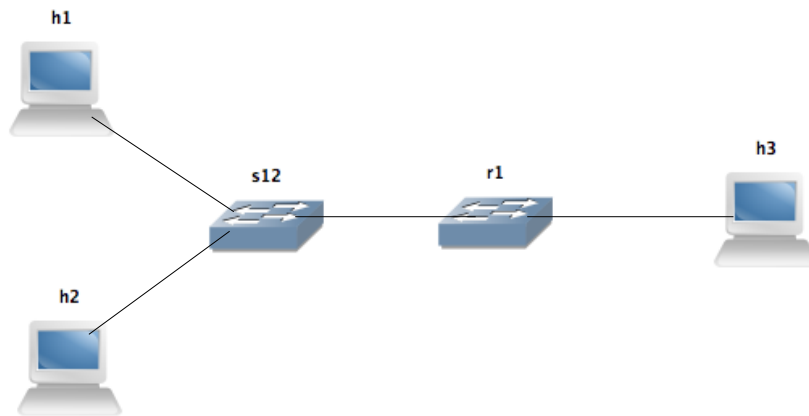


Figure 2: Network configuration with a router and a switch

```
mininet> exit
# mn -c
# python topo2.py
```

Perform a reachability test in Mininet. A reachability test is a test to determine which hosts can 'reach' one another. This is performed by having each host ping all other hosts. In our case, this also includes the router. A quick way to do this test in Mininet is by running the following command.

```
mininet> pingall
```



Q13/ What is the percentage of dropped packets? What does this correspond to?

[A13]

Q14/ Which hosts are unable to reach one another?

[A14]

We will now attempt to fix the problem. First, open the `topo2.py` script and inspect it.



Q15/ What is the subnet mask used throughout the file?

[A15]

Q16/ What are the interfaces and respective IP addresses of the router `r1`?

[A16]

Q17/ Can you spot any misconfigurations in the file?

[A17]

The first issue we notice in the file is that IP forwarding has been disabled in the router `r1`. We will enable it and attempt the reachability test again.



Q18/ What is the command you will use to enable IPv4 forwarding on `r1`?

[A18]

Q19/ What is the percentage of dropped packets now in the reachability test? Which hosts are still unable to ping each other?

[A19]

The next issue in the file is the default gateway of `h2`, which is set to `10.0.0.101`. This IP address does not exist in our network. Delete this configuration and set a proper default gateway through the terminal of `h2`.



Q20/ What are the commands you used to achieve this?

[A20]

Q21/ What are the results of the reachability test now?

[A21]

Now we learn about routing between two interfaces. Monitor the traffic in Wireshark of both interfaces of `r1`. From `h1`, ping `h3`.



Q22/ What changes are done to IP packets when they are routed between `r1`'s `r1-eth0` and `r1-eth1`?

[A22]

Q23/ What is the purpose of such changes?

[A23]

2.4 ROUTING WITH MULTIPLE HOPS

In this section we want to see what happens when we introduce a second intermediate routing device to the network. Let's start by setting the default gateway of `r1` to `h3`'s IP address, then on `h3` remove the default gateway and enable IPv4 forwarding.



Q24/ Type in the commands you need to do this.

[A24]

Monitor `eth0` of `h3` and `r1`. Try pinging from `h1` to `h3` and `h2` to `h3`.



Q25/ Based on Wireshark captures, explain why it does not work.

[A25]



Q26/ Write down **one single command** that fixes the problem (namely pinging from `h1` and `h2` to `h3`), and specify the PC where you need to apply it.

[A26]

Ping again from `h1` and `h2` to `h3`, and confirm that your fix solves the problem before moving to the next section.

3 CONNECTING VIRTUAL ENVIRONMENT TO THE REAL WORLD USING NETWORK ADDRESS TRANSLATION (NAT)

In this section we will use what we learned from Lab1 about manipulating the `iptables` filter. The purpose of the section is to connect the isolated virtual network that we have deployed so far, to the real Internet.

We will work in the network described in Figure 3. `h1` and `h2` are workstations, `r1` is an aggregation router, and `h3` is the perimeter router where we will have our connection to the real world.

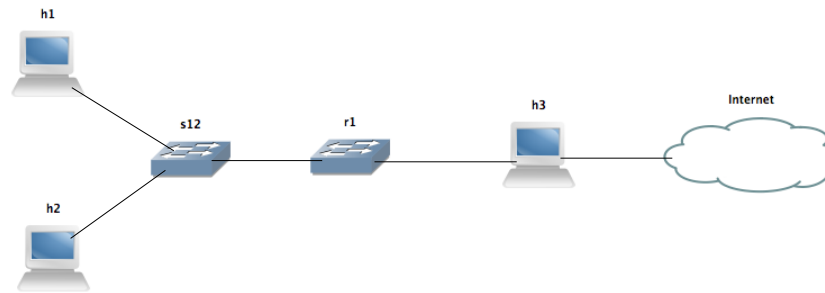


Figure 3: Network configuration with a connection to the real world



Q27/ We have one real connection (IP address) to the real world, but we have two clients (`h1` and `h2`) that require access to the internet. Which solution would you use to tackle this problem, and explain how would it solve the problem.

[A27]

There are two main steps to connect your virtual environment to the real Internet:

1. We require a real IP address on `h3-eth1` interface of `h3`.
2. We need to masquerade the traffic coming from `h1` and `h2`.

3.1 BORROWING AN IPV4 ADDRESS

As we said before, we need an *existing* (real) IPv4 address that can be used to connect to the Internet. The purpose of this section is to obtain such valid IPv4 address. There are three steps for that:

1. Create a bridge between a real interface in your host machine, and `h3's eth1`.
2. Finding a suitable IPv4 address.
3. Setting up the IPv4 address of `h3-eth1` on `h3`.

3.1.1 BRIDGE BETWEEN THE PHYSICAL AND THE VIRTUAL INTERFACE

For the first step, the process is shown schematically in Figure 4, and we will cover it step by step.

To perform the bridging between physical and virtual network-adapters, we first need to find out which interface our VM is using to access the internet.

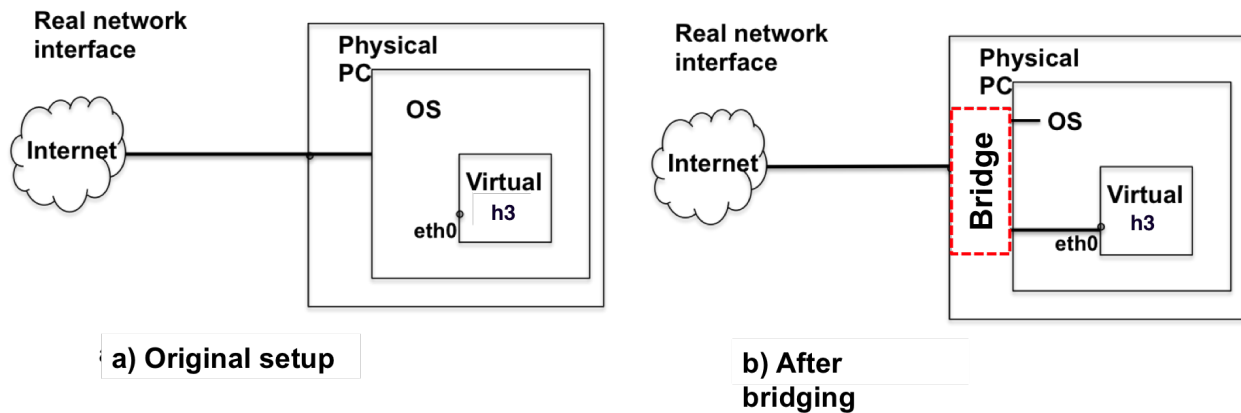


Figure 4: Bridging the network adapter

Exit Mininet, clean up the topology, and run wireshark on the VM. Start capturing on each interface and ping `www.google.com` to find out if that interface is being used to access the internet.

```
mininet> exit
# mn -c
# wireshark &
# ping www.google.com
```



Q28/ What is the IP address of the interface used to access the internet? Do you notice anything special about it?

[A28]

Now that we know which interface is used by the VM, the next step is to implement a bridge between it and `h3-eth1`.

In Mininet, we can only implement a bridge with an OVS. To mitigate this, we will add an OVS between `h3` and the internet cloud in Figure 3.

Run the script `topo2_internet.py`. This will create the same topology as before (in addition to all the extra configurations that you have performed, and addition of an extra OVS that we will use to perform the bridge).

From Mininet, execute the following command.

```
mininet> sh ovs-vsctl add-port s3 <interfacename>
```

Replace `interfacename` with the interface that accesses the Internet in your VM. Remember that the sudo password for the VMs is `lca2`. Now head over to the terminal of `h3` and run the following command.

```
# dhclient h3-eth1
```

This automatically sets a usable IP address to the `eth1` interface of `h3`, allowing it to access the internet through the bridge we just set up. Test the configuration by pinging Google or EPFL from `h3`.

3.1.2 CHOOSING THE BEST IP ADDRESS FOR `h3`'S `ETH1`

The task in this section is to find a suitable IP address for `h3`'s `eth1`. This was basically achieved using the `dhclient` command in the previous section. Here, we will learn to choose an appropriate IP address manually.

Run `ifconfig` and `route -n get default` (in Linux or Mac) or `ipconfig /all` (in Windows) and check your physical IP address, subnet mask and default gateway.



Q29/ Write down your physical IP address, subnet mask and default gateway. According to your network configuration, what is the range of IP addresses that we could use for the virtual adapter?

[A29]

Q30/ How could you find the current non-used IP addresses in your LAN?. Is it safe to take any of them?

[A30]



Q31/ If you are doing this exercise outside EPFL (e.g. at your home), you will most likely get a private IP address and default gateway (e.g. `192.168.x.x` or `10.x.x.x`). Can you use these private IP addresses on `h3`'s `eth1` to do NAT? Explain why.

[A31]

3.2 NAT CONFIGURATION

We worked with the command used to configure NAT in `Lab1`, `iptables -t nat`, which manages the table that contains rules regarding address translations. In this section, we will analyze how NAT works for different types of packets.

First let's see what happens when `h1/h2` access the Internet with their native IP address. Monitor with Wireshark the interface `eth1` of `h3`. From `h1/h2`, ping to Google and its IP address.

```
# ping -c 5 www.google.com
```

```
# ping -c 5 172.217.18.100
```



Q32/ Analyze the packets coming from h1/h2 and explain why you are unable to reach Google.

[A32]

Q33/ Propose the `iptables -t nat` command you need to properly configure NAT in h3.

[A33]

Test from h1 and h2 and you have Internet connectivity by pinging Google and test that you have successfully configured your router to do NAT.

Next, let's explore how NAT works!!.

Do `traceroute` to Google from h2 and then from h3, while capturing `eth0` and `eth1` traffic on h3 using Wireshark. Explore the difference in the traffic on both cases.



Q34/ When doing `traceroute` from h2, what is the difference in the packets captured on h3's `eth0` and `eth1`?

[A34]



Q35/ Focus on the `traceroute` from h3. What is the difference in the packets as compared to h2?

[A35]



Q36/ Which field in the UDP packet is used to identify the (local) source IP address of h2 in order to properly forward incoming ICMP replies back to it?

[A36]

Do `ping` to Google from h1 and h3, while capturing the traffic on h3 (both on `eth0` and `eth1`) using Wireshark. Explore the difference in the traffic in both cases.



Q37/ What is the difference in the request ICMP packets captured between packets sent from h1 and packets sent from h3 when capturing on the exit interface of each?

[A37]



Q38/ Conclude how the incoming ICMP replies are forwarded back to h1 when doing ping from h1. In particular, which field in the request/reply ICMP packets was used to identify the (local) source IP address?

[A38]

4 HAVING INTERNET ACCESS THROUGH A MACHINE IN INTERNET ENGINEERING WORKSHOP (IEW)/INF019

Up to now, you were accessing the Internet by directly connecting your laptop to a WiFi access point (AP) or Ethernet socket at your home or at EPFL. In this section, you will access the Internet by connecting your laptop to a computer in INF019 using an Ethernet cable. The goal of this section is to give you a feel about the real communication networks. The communication network in IEW is very simple and small, but is probably enough to give you a glance of real network devices and networking between them. You can find the schema of the communication network in IEW [here](#). Also, please watch the [video](#) before you start working on the Minix machines (and the virtual machines on them) in the IEW (INF 019).

To start with, please disable the WiFi interface of your laptop and plug one end of an Ethernet cable into your laptop and the other in one of the available interfaces of a computer in INF019. Now, configure the connected network interfaces of these two machines (your laptop and the computer from INF019) such that you can ping the lab computer from your laptop and your laptop from the lab computer. To be able to ping these machines with each other, you will first need to assign ip addresses to these two connected interfaces. Please choose a subnetwork (in private ipv4 address space) that is not yet used in this setting and assign two different ip addresses to these machines.



Q39/ To be able to ping between your laptop and the computer in the lab, please write all the commands you executed on your laptop?

[A39]



Q40/ To be able to ping between your laptop and the computer in the lab, please write all the commands you executed on the computer in the lab?

[A40]

Now, as you are able to ping these two machines, the next goal is to be able to ping a machine on the Internet. For this, we will use the other interface of the machine in the lab, which is connected to the Internet. This means, your machine will be connected to the machine in the lab, which, in turn, is connected to the Internet with the other interface.

Now, configure the lab computer and your computer such that you can ping Google DNS server 8.8.8.8.



Q41/ To do the above, please write all the commands you executed on the lab computer?

[A41]



Q42/ To do the above, please write all the commands you executed on your laptop?

[A42]

Once you successfully tested pinging to 8.8.8.8 ip address from your laptop, can you, try pinging `google.com`? If you have not yet specified an address of a DNS server for your laptop, the next thing is to do that. Once done, please, try to ping `google.com` and you should now be able to ping it.

To find the ip address of a DNS server, you can either use the DNS server address used by the lab computer (on the Interface that is connected to the Internet) or you can use Google's public DNS server (8.8.8.8).



Q43/ How did you configure the DNS server on your laptop?

[A43]

If everything is well configured, you should now be able to access the Internet from your laptop connecting through the lab computer. Please visit a web server (`facebook.com`, `google.com`) from your browser to confirm it.

In section 3, you have already connected your mininet hosts to the Internet. Follow the steps from section 3 or if you have saved the commands in a script for section 3, execute them so that you can have the same topology as you had at the end of section 3. You should now be able to access the Internet from hosts h1, h2, and h3.



Q44/ Do the `traceroute` for `epfl.com` from mininet host `h1` and paste the results below.

[A44]

5 TROUBLESHOOTING

As the title for this section suggests, in networking things will not always work out as expected. Before starting to work in this section, you will have to execute some executables. These executables will put a number of PCs into a problematic situation where something doesn't work. Your task is to find out what the problem is and propose a solution.

You should not perform any debug command or Wireshark capture in h3. Assume h3 is a router controlled by your Internet service provider (ISP) and you don't have access to it.

5.1 ABOUT GRADING THIS SECTION

The points given for your answer mostly depend on your explanations about how you located the problem!, so describe precisely your steps to locate and diagnose the fault. You should use the scientific method when answering the problem. The methodology you should use is the following:

1. Pose a hypothesis
2. Run experiments to validate the hypothesis
3. If validation is OK exit, else loop (go back to 1. by posing another hypothesis)

In your answer you should write down all steps. Specially, you should also write down all hypotheses that later proved to be wrong. We want to see the path you took to reach your final conclusion!

More specifically for this Lab: What were the commands you executed to get there? Up to which point did the system work as expected? What were the actions that never got executed but were expected? What was the packet that did not reach its destination? Where and why did it get dropped/lost?

5.2 WE WERE HACKED!!!

Jon and Arya are roommates and close friends. Both of them are connected to the Internet through the same home router thus the same ISP. The configuration is the same as described in Figure 5, where h1 is Jon's computer and h2 is Arya's computer, r1 is the home router, shared between Jon and Arya, and h3 is the ISP's router.

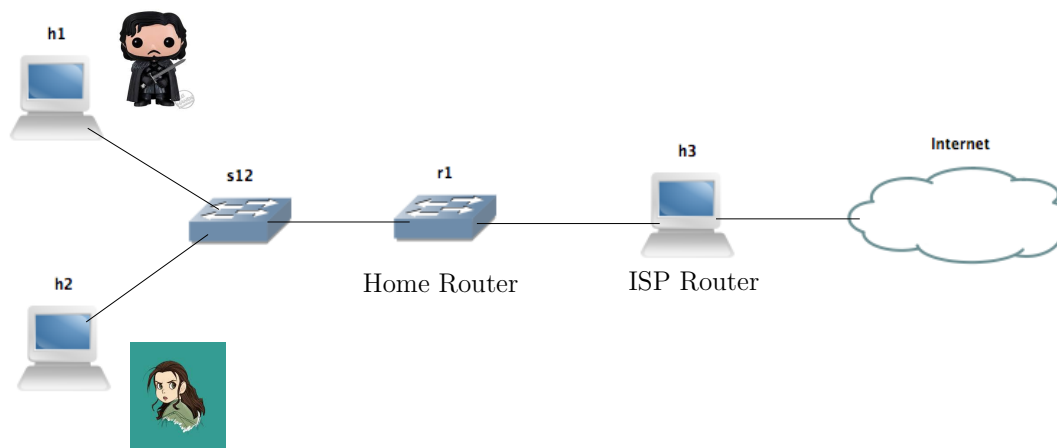


Figure 5: Troubleshooting configuration

Jon and Arya are foreign students at EPFL and they use Facebook a lot to communicate with their relatives in their home town, Winterfell. None of them is an expert in computers and networking. They have a common enemy, Joffrey, who plays jokes and is a Computer Science student and network expert. Joffrey told Jon and Arya that he hacked their computers (h1 and h2) as well their home router (r1) and that they would never navigate through Facebook again.

To simulate Joffrey's malicious attack, we prepared an executable for h1, h2 and r1. You should download the executables from the course web page on moodle. Unzip and extract the corresponding files:

```
# unzip lab2-execs.zip
```

We recommend you to extract all executables to your shared folder between host and virtual machines.

Let's start with Jon and h1. Check that you are on host h1. Navigate to the lab2-execs folder and execute the executable for h1 on h1. Make sure you **run the executable as a superuser in Linux**.

```
# cd lab2-execs
# sudo ./h1.hack
```

Note that if you stop the simulation in Mininet, then you need to redo the previous section and make sure that h1 and h2 have a proper internet connection, then reload the script on h1.

Also, if you do not manage to solve one of the questions, you may want to restart Mininet, redo the previous section and make sure that all hosts have an internet connection, and then move on to the next question. This is where using scripts to save the commands you performed to achieve a working configuration would help a great deal. Ideally, solving a question means undoing the changes brought about by the hack, so you can start solving any of the questions in this section from the stable working configuration of Section 3.

Now, your mission is to help Jon find out why he cannot navigate to Facebook anymore. You are only allowed to have Wireshark captures and use debug commands in h1.

Open firefox on h1 by typing `firefox &` in the corresponding terminal.



Q45/ Are there any problems in Jon's PC (h1)? Enumerate them (if any) and write down how you find them and how you fix them.

[A45]

After you fix the issue, confirm that Jon can now safely log into Facebook to update his status to "I know nothing about TCP/IP".

Now repeat the same procedure for Arya: execute the `h2_hack` executable on the `h2` terminal. Again you can use Wireshark on `h2` to capture packets.



Q46/ Open Firefox on `h2` and try to navigate to Facebook. Are there any problems in Arya's PC (`h2`)? Enumerate them (if any) and write down how you find them and how you fix them.

[A46]

While you were solving these issues on `h1` and `h2`, Joffrey had enough time to hack `r1`. To simulate this, execute the `r1_hack` executable on `r1`. Doing this, Joffrey managed to stop both Jon and Arya from accessing Facebook again. Your job is to get both of them Facebook access again. Hurry though, Winter is Coming!

For this question, you may perform debugging and capture packets using Wireshark on `h1`, `h2`, and `r1`. DO NOT TOUCH `h3`!



Q47/ Enumerate the problems (if any) that were stopping Jon (`h1`) from accessing Facebook. How did you find them and fix them?

[A47]

Q48/ Enumerate the problems (if any) that were stopping Arya (`h2`) from accessing Facebook. How did you find them and fix them?

[A48]